

Lab 4 : Kubernetes DaemonSets

Nom du module : PaaS Environment

Année Universitaire : 2021/2022

Public cible : 4 ArcTIC

Enseignants :

- Soumaya Mbarek
- Hamdi Gabsi

Goals :

In this lab, we will have a closer look to Kubernetes DaemonSets. Daemonsets are similar to Kubernetes Deployments, but they will deploy one and only one single POD per (non-master) Kubernetes node. We will have a closer look at the two possible update strategies of DaemonSets: RollingUpdate (default) and OnDelete.

I. Step 0 : Access the Kubernetes Playground

We start by accessing the [Katacode Kubernetes Playground](#).

II. Step 1: Create a DaemonSet

Let us try to create a Daemonset using the following YAML file :

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: frontend-ds
spec:
  selector:
    matchLabels:
      tier: my-matching-label
  template:
    metadata:
      labels:
        tier: my-matching-label
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```

III. Step 2: Change the Update Strategy to OnDelete

Let us review the update strategy of the DaemonSet as follows:

Lab 4 : Kubernetes DaemonSets

```
controlplane $ kubectl get daemonset -o yaml | grep -A 2 updateStrategy
  f:updateStrategy:
    f:rollingUpdate:
      .: {}
--
updateStrategy:
  rollingUpdate:
    maxUnavailable: 1
```

We can see, that the default update strategy of a DeamonSet is a RollingUpdate. Let us change that to OnDelete for now.

```
Terminal Host 1 +
  imagePullPolicy: IfNotPresent
  name: php-redis
  resources: {}
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
  dnsPolicy: ClusterFirst
  restartPolicy: Always
  schedulerName: default-scheduler
  securityContext: {}
  terminationGracePeriodSeconds: 30
  updateStrategy:
    rollingUpdate:
      maxUnavailable: 1
    type: OnDelete
```

IV. Step 3: Change POD Template Image Version

We now edit the file again and change the POD template image version from v3 to v2:

```
Terminal Host 1 +
  apiVersion: apps/v1
  kind: DaemonSet
  metadata:
    name: frontend-ds
  spec:
    selector:
      matchLabels:
        tier: my-matching-label
    template:
      metadata:
        labels:
          tier: my-matching-label
      spec:
        containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v2
```

Lab 4 : Kubernetes DaemonSets

We check the POD Image version(s):

```
Terminal Host 1 +
controlplane $ kubectl describe pods | grep image
Normal Pulled 26m kubelet, node01 Container image "gcr.io/google_sam
ples/gb-frontend:v3" already present on machine
```

We can see, that the image version has not changed. To trigger a replacement, we need to delete the pod:

```
controlplane $ kubectl get pods
NAME READY STATUS RESTARTS AGE
frontend-ds-q6x26 1/1 Running 0 27m
controlplane $ kubectl delete pod frontend-ds-q6x26
pod "frontend-ds-q6x26" deleted
```

Now, we expect a new POD to be created with the new Docker image:

```
controlplane $ kubectl get pods
NAME READY STATUS RESTARTS AGE
frontend-ds-rhvm9 1/1 Running 0 2m40s
controlplane $ kubectl describe pod frontend-ds-rhvm9 | grep image
Normal Pulling 3m29s kubelet, node01 Pulling image "gcr.io/google_samp
les/gb-frontend:v2"
Normal Pulled 2m42s kubelet, node01 Successfully pulled image "gcr.io
/google samples/gb-frontend:v2"
```

V. Step 4: Rollout and Rollback

1. Step 4.1: Rollback History

We can see that there is a rollout history available similar to Deployments:

```
controlplane $ kubectl rollout history daemonset frontend-ds
daemonset.apps/frontend-ds
REVISION CHANGE-CAUSE
1 <none>
2 <none>
```

2. Step 4.2: Rollback

We can rollback to Image version 3 by undoing the previous rollout. This will roll back the image version of the POD template to v3. However, the PODs will not be updated automatically, and we need to delete the POD to force the update.

Lab 4 : Kubernetes DaemonSets

```
controlplane $ kubectl rollout undo ds
daemonset deployment statefulset
controlplane $ kubectl rollout undo ds frontend-ds
daemonset.apps/frontend-ds rolled back
controlplane $ kubectl describe pods frontend-ds-2vnbq | grep image
Normal Pulling 2m56s kubelet, node01 Pulling image "gcr.io/google_samp
les/gb-frontend:v2"
Normal Pulled 2m37s kubelet, node01 Successfully pulled image "gcr.io
/google_samples/gb-frontend:v2"
controlplane $ kubectl delete pod frontend-ds-2vnbq
pod "frontend-ds-2vnbq" deleted
```

```
Terminal Host 1 +
controlplane $ kubectl describe pod frontend-ds-njgwt | grep image
Normal Pulled 73s kubelet, node01 Container image "gcr.io/google_sam
ples/gb-frontend:v3" already present on machine
```

VI. Step 5: Rolling Update

Now let us change the Update Strategy to RollingUpdate:

```
Terminal Host 1 +
resources: {}
terminationMessagePath: /dev/termination-log
terminationMessagePolicy: File
dnsPolicy: ClusterFirst
restartPolicy: Always
schedulerName: default-scheduler
securityContext: {}
terminationGracePeriodSeconds: 30
updateStrategy:
  rollingUpdate:
    maxUnavailable: 1
    type: RollingUpdate
status:
  currentNumberScheduled: 1
  desiredNumberScheduled: 1
```

Now let us change the version again. This time, we choose v1:

```
Terminal Host 1 +
tier: my-matching-label
template:
  metadata:
    creationTimestamp: null
    labels:
      tier: my-matching-label
  spec:
    containers:
      - image: gcr.io/google_samples/gb-frontend:v1
        imagePullPolicy: IfNotPresent
        name: php-redis
        resources: {}
        terminationMessagePath: /dev/termination-log
        terminationMessagePolicy: File
    dnsPolicy: ClusterFirst
    restartPolicy: Always
```

Lab 4 : Kubernetes DaemonSets

Since the update strategy is RollingUpdate now, the POD(s) are replaced immediately:

```
controlplane $ kubectl edit ds frontend-ds
daemonset.apps/frontend-ds edited
controlplane $ kubectl describe pod frontend-ds-jpx4p | grep image
Normal Pulling      26s  kubelet, node01    Pulling image "gcr.io/google_sampl
es/gb-frontend:v1"
Normal Pulled       12s  kubelet, node01    Successfully pulled image "gcr.io/
google_samples/gb-frontend:v1"
```

Note : We have created a minimal Kubernetes DaemonSet. First, we have observed that and POD template changes are not propagated to existing PODs if we choose the OnDelete update strategy. In that case, PODs need to be deleted, to trigger the update of the PODs. However, if we choose the RollingUpdate strategy, POD renewal is triggered with any update of the DaemonSet's POD template.